

Sparse Recovery and Compressed Sensing

A Model-Based Deep Learning Approach

Core question: How far can model-based structure take us — and when do we need data-driven flexibility?

Etay Baron • Etai Wigman

Model-Based Deep Learning (361.2.2320) • Ben-Gurion University • Spring 2025–2026

Based on: Chen, Liu, Wang, Yin — "Hyperparameter Tuning is All You Need for LISTA" • NeurIPS 2021

Seminar Roadmap

1

Background

Motivation, LASSO,
ISTA / FISTA

2

Unrolling

LISTA, ALISTA,
HyperLISTA

3

Our Study

Scalar variants +
training methods

4

Results

Sparse recovery &
design space

5

Images

MNIST & FashionMNIST
CS stress tests

Goal: systematically compare how much learning is needed when moving from ideal sparse signals to real compressed-sensing images.

Synthetic signals

$m=250$, $n=500$, $s=50$
Gaussian A , noiseless

Trained models

LISTA, ALISTA, HyperLISTA
+ 3 scalar designs (ours)

Image CS

MNIST & FashionMNIST pixel domain
ratios $\in \{0.25, 0.5\}$

Motivation & Problem Formulation

Sparse recovery appears everywhere:

Application	Measurements	Goal
MRI acceleration	k-space ($m \ll n$)	Image pixels
Radar	Compressed snapshots	Reflectivity
Hyperspectral	Spectral bands	Spatial cube
Compressed sensing	Random projections	Original signal

Why tractable? Natural signals have only $s \ll n$ non-zero components.

LASSO — the optimization backbone:

$$\hat{x} = \arg \min_x \frac{1}{2} \|b - Ax\|_2^2 + \lambda \|x\|_1$$

Measurement model

$$b = Ax^* + \varepsilon$$

$A \in \mathbb{R}^{m \times n}$, $m \ll n$

x^* : s -sparse signal

ε : additive Gaussian noise

Metrics: NMSE (dB) for synthetic signals;
NMSE + PSNR + SSIM for image CS.

Classical Baselines: ISTA and FISTA

ISTA — proximal gradient descent on the LASSO objective:

$$x^{(k+1)} = S_{\theta} \left(x^{(k)} + \frac{1}{L} A^T (b - Ax^{(k)}) \right), \quad L = \|A\|_2^2, \quad \theta = \frac{\lambda}{L}$$

FISTA — ISTA with Nesterov momentum, $O(1/k^2)$ convergence vs $O(1/k)$:

$$y^{(k)} = x^{(k)} + \frac{t_k - 1}{t_{k+1}} (x^{(k)} - x^{(k-1)}), \quad x^{(k+1)} = S_{\theta} \left(y^{(k)} + \frac{1}{L} A^T (b - Ay^{(k)}) \right)$$

Method	NMSE @ K=16	Params
ISTA	-5.3 dB	0
FISTA	-11.0 dB	0

At K=16 iterations both methods are far from converged.

ISTA needs 200+ iterations for -20 dB.

Can we make 16 steps count?

Deep Unrolling: From Iterations to Layers

Key idea (Gregor & LeCun, 2010): unfold K ISTA iterations into a K-layer network and learn the update parameters from data.

Classical ISTA — parameters fixed

$$z^{(k)} = x^{(k)} + \frac{1}{L} A^T (b - Ax^{(k)})$$
$$x^{(k+1)} = S_{\theta}(z^{(k)})$$

Parameters $1/L$ and $\theta = \lambda/L$ are fixed constants.

Fixed — no learning

Unrolled LISTA — parameters learned

$$z^{(k)} = W_x^{(k)} x^{(k)} + W_y^{(k)} b$$
$$x^{(k+1)} = S_{\theta_k}(z^{(k)})$$

$W_x^{(k)}, W_y^{(k)}, \theta_k$ are learned per layer via BPTT.

~6 M parameters — learned end-to-end

MBDL principle: preserve the known algorithmic structure (gradient step + soft-threshold);
learn only what improves finite-depth reconstruction. Training minimizes $\|x(K) - x^*\|^2$ via BPTT.

The MBDL Ladder: LISTA → ALISTA → HyperLISTA

LISTA — learn all matrices per layer, ~6 M params:

$$x^{(k+1)} = S_{\theta_k} (W_x^{(k)} x^{(k)} + W_y^{(k)} b)$$

ALISTA — W analytic, learn step sizes and thresholds, 32 params:

$$W^* = (AA^T)^{-1}A \text{ (normalized)}$$

$$x^{(k+1)} = S_{\theta_k} (x^{(k)} + \gamma_k W^{*T} (b - Ax^{(k)}))$$

HyperLISTA — adaptive $\theta/\beta/p$ from residuals, 3 params (grid search, zero backprop):

$$\theta^{(k)} = c_1 \mu \|A + (Ax^{(k)} - b)\|_1, \quad \beta^{(k)} = c_2 \mu \|x^{(k)}\|_0$$

$$p^{(k)} = c_3 \log \frac{\|A + b\|_1}{\|A + (Ax^{(k)} - b)\|_1}$$

Model	Params	NMSE
ISTA	0	-5.3 dB
FISTA	0	-11.0 dB
LISTA	6 M	-20.3 dB
ALISTA	32	-29.8 dB
HyperLISTA	3	-61.4 dB

More structure → fewer params → better performance

This is the MBDL principle in action.

Our Design Space Study: What Should We Learn?

Research question: LISTA learns 6 M parameters; HyperLISTA uses 3 scalars. What is the minimal useful set of learnable components?

The ISTA gradient-correction update has two free scalar sequences:

$$x^{(k+1)} = S_{\theta_k}(x^{(k)} + \gamma_k A^T(b - Ax^{(k)}))$$

We designed three models, each learning a different subset:

ThresholdISTA

16 params

Learns $\theta_1, \dots, \theta_K$ only

Step fixed at $\gamma = 1/L$

Init: $\theta_0 = \lambda/L$

StepISTA

16 params

Learns $\gamma_1, \dots, \gamma_K$ only

Threshold coupled: $\theta_k = \lambda\gamma_k$

Init: $\gamma_0 = 1/L$

StepThresholdISTA

32 params


Learns both γ_k and θ_k

independently per layer

Init: ISTA values

All parameters softplus-reparameterized (positivity). All initialized to reproduce ISTA exactly before training.

Our Results: Scalar Models



Model	NMSE	Params
ISTA	-5.3 dB	0
FISTA	-11.0 dB	0
ThresholdISTA	-7.1 dB	16
StepISTA	-20.0 dB	16
StepThresholdISTA	-23.6 dB	32
LISTA-L1 (ref.)	-20.3 dB	6 M

Key finding

Learning γ_k alone (16 scalars) already matches full LISTA (6 M parameters).

Decoupling θ_k adds 3 dB at zero extra cost.

StepThresholdISTA beats LISTA-L1 with 187,000× fewer parameters.

Training Methods for Unfolded Optimizers

Architecture fixed (LISTA). Only the training strategy changes.

L1 — End-to-end

$$\mathcal{L} = \|x^{(K)} - x^*\|^2$$

Supervise final output only.
Standard BPTT.

L2 — Deep Supervision

$$\mathcal{L} = \|x^{(K)} - x^*\|^2 + \lambda \sum_k w_k \|x^{(k)} - x^*\|^2$$

Supervise every
intermediate layer.

L3 — Greedy Layerwise

$$\min_{\theta_k} \|x^{(k)} - x^*\|^2$$

Train layer k while
layers 1...k-1 frozen.

Tied vs. Untied weights:

	Untied LISTA	Tied LISTA
Weights	Independent per layer	Single shared (Wx, Wy, θ)
Params	~6 M	375 K
Style	Fully flexible	RNN-like
Effect	Hard to optimize	Smoother loss landscape


Weight tying = implicit regularization.

LISTA-Tied-L1: -19.1 dB @ 375 K

LISTA-L1: -20.3 dB @ 6 M

Synthetic Sparse Recovery Results

Main unfolded optimizers: NMSE vs layer




Method	NMSE	Params
ISTA	-5.3 dB	0
FISTA	-11.0 dB	0
LISTA-L1	-20.3 dB	6 M
LISTA-L2	-21.9 dB	6 M
LISTA-L3	-16.6 dB	6 M
LISTA-Tied-L1	-19.1 dB	375 K
ALISTA	-29.8 dB	32
HyperLISTA	-61.4 dB	3

L3 plateau: NMSE flattens ~layer 8.
Greedy training gives no incentive to improve beyond what prior layers solved.

LISTA-Tied > LISTA-Untied despite 16× fewer params.

Design-Space Map



Model	Params	NMSE	What drives it
ThresholdISTA	16	-7.1 dB	learned θ_k
StepISTA	16	-20.0 dB	learned γ_k
StepThresholdISTA	32	-23.6 dB	$\gamma_k + \theta_k$
LISTA-L1	6 M	-20.3 dB	full matrices
ALISTA	32	-29.8 dB	analytic W
HyperLISTA	3	-61.4 dB	adaptive residuals

StepThresholdISTA (32 scalars)
beats LISTA-L1 (6 million matrices).
The step-size schedule γ_k is the bottleneck.

Real Image CS: FashionMNIST Pixel Domain

Setup: flatten 28×28 image to $x \in [0,1]^{784}$, measure $b = Ax$, ratio $m/784 \in \{0.25, 0.5\}$. FashionMNIST: ~51.5% of pixels exactly zero.

Method	NMSE	PSNR	SSIM	Params
ISTA	−1.2 dB	8.9	0.13	0
FISTA	−1.2 dB	8.9	0.14	0
StepThresholdISTA	−1.2 dB	8.9	0.13	32
ALISTA	−0.9 dB	8.4	0.12	32
HyperLISTA	−0.3 dB	7.7	0.09	3
LISTA-L2	−13.7 dB	22.3	0.82	12 M

ratio = 0.25

Signal model mismatch

All structured methods assume i.i.d.

Gaussian sparse signals.

FashionMNIST pixels are bounded $[0,1]$, spatially correlated, and only soft-sparse.



Why LISTA wins

LISTA-L2 is assumption-free —
it learns the actual image distribution
from 51 K training images.

Wrong prior is worse than no prior.

Robustness Across Datasets and Measurement Ratios

SSIM across {MNIST, FashionMNIST} × ratio {0.25, 0.5}: harder data and fewer measurements widen the gap.





Trend

MNIST @ 0.5 is easy for everyone; FashionMNIST @ 0.25 breaks every structured model. LISTA-L2 stays ≥ 0.82 SSIM everywhere.

Interpretation

The sparse-Gaussian prior degrades gracefully on MNIST (~78% zeros) but fails on FashionMNIST (~51% zeros, textured).

Visual Stress Tests (ratio = 0.25)



LISTA-L2 recovers clean structure on both datasets. ALISTA / HyperLISTA zero out valid signal on FashionMNIST — the Gaussian-sparse threshold is miscalibrated for textured images. Wrong prior is worse than no prior.

Conclusions

1 Structure wins when assumptions match.

HyperLISTA's 3 scalars beat 6 M learned parameters by 41 dB on synthetic sparse data.

2 Step sizes matter most.

StepISTA (16 scalars) matches LISTA. Decoupling θ_k adds 3 dB. Matrix learning is largely redundant.

3 Weight tying is implicit regularization.

LISTA-Tied matches LISTA-Untied (-19.1 vs -20.3 dB) with 16× fewer parameters.

4 Wrong prior actively hurts.

ALISTA / HyperLISTA score below ISTA on real images. LISTA wins by making no structural assumptions.

MBDL hierarchy on synthetic data:

ISTA → LISTA → ALISTA → HyperLISTA

0 → 6M → 32 → 3 params

-5 → -20 → -30 → -61 dB

Future work

Apply models in DCT / wavelet domain — images are sparser there, closer to the Gaussian-sparse model.

Design image-aware thresholds for bounded, non-negative pixel structure.

Code & References

Code — public GitHub repository (notebooks 01-04 + src modules):

github.com/ThEpiCake/hyperlista-mbd

References:

- [1] A. Beck and M. Teboulle, "A Fast Iterative Shrinkage-Thresholding Algorithm for Linear Inverse Problems," SIAM J. Imaging Sciences, 2009.
- [2] K. Gregor and Y. LeCun, "Learning Fast Approximations of Sparse Coding," ICML 2010.
- [3] X. Chen, J. Liu, Z. Wang, W. Yin, "Theoretical Linear Convergence of Unfolded ISTA and Its Practical Weights and Thresholds," NeurIPS 2018.
- [4] J. Liu, X. Chen, Z. Wang, W. Yin, "ALISTA: Analytic Weights Are As Good As Learned Weights in LISTA," ICLR 2019.
- [5] X. Chen, J. Liu, Z. Wang, W. Yin, "Hyperparameter Tuning is All You Need for LISTA," NeurIPS 2021.
- [6] N. Shlezinger, J. Whang, Y. C. Eldar, A. G. Dimakis, "Model-Based Deep Learning," Proceedings of the IEEE, 2023.
- [7] Course lecture notes 1-8, Model-Based Deep Learning (361.2.2320), Ben-Gurion University, 2026.

Thank you! Questions welcome.